

# *An Introduction to Cluster Computing*

**With ordinary things, do extraordinary things**

Soma

Disclaimer:

The contents are not validated and the reliability is not checked However, efforts were put ensure that they are true to best of my knowledge

## What is a cluster?

- A group of connected workstations/ computers that can handle the job/ task in a shared manner.
  - $1 + 1 = 3$

## What is cluster Computing?

- It represents an attempt to solve larger problems, or to solve problems in a more cost effective manner, than the more conventional systems of the time.

***Speed – Due to Compute-bound problem and/ or I/O-bound problem, consider parallelism***

# What is parallelism?

- It all depends on the application
- CONCURRENT parts of the algorithm are those that can be computed independently
- PARALLEL parts of the program are those CONCURRENT parts that are executed on separate processing at the same time
- Computation & Implementation are **NOT** interchangeable
- Cluster Computing is a way to exploit (forced) CONCURRENCY!

***Algorithm is not parallel. Hardware makes the concurrent parts of the algorithm RUN parallel***

# Why a Cluster/ Why not a Cluster ?

- Performance of Commodity & network hardware increase and their prices decrease, it is more practical to build parallel computers
- Buying single large CPU time: [History](#)  
Group off-the shelf commodities to build parallel computers:  
[Present](#)  
Home Super Computers: [Future](#)
- CPU speeds have been **DOUBLING** every 18 months  
Processors speeds will **NOT** continue to do so after 2005
- It is possible to do several things at once with MTOS. This is natural parallelism that can be exploited by clusters

***Parallel computing speeds up things anywhere from 2 to 500***

# Platforms for Computing

- Single-tasking OS: MS-DOS
- Multi-tasking OS: UNIX, NT using single CPU
- Multi-tasking OS with Multiple CPUs: UNIX, NT with more than one CPU
- Threads on Multi-tasking OS with Multiple CPUs:
- UNIX, NT with more than one CPU on the same motherboard

***Many ways to implement parallel computing. Consider portability, scalability, performance, and COST before making the decision***

## **Hardware Architecture:**

Local Memory Machines that communicate through messages (Beowulf)

Shared Memory machines that communicate through memory (COW)

## **Software Architecture:**

Using Messages sent between processors

Using Operating system threads

Locomotives (olden days) were connected in

- Shunt when you need speed
- Series when you want starting torque

***A major classification of cluster computing  
Cluster of Workstations (COWs) Vs Beowulf***

## Beowulf

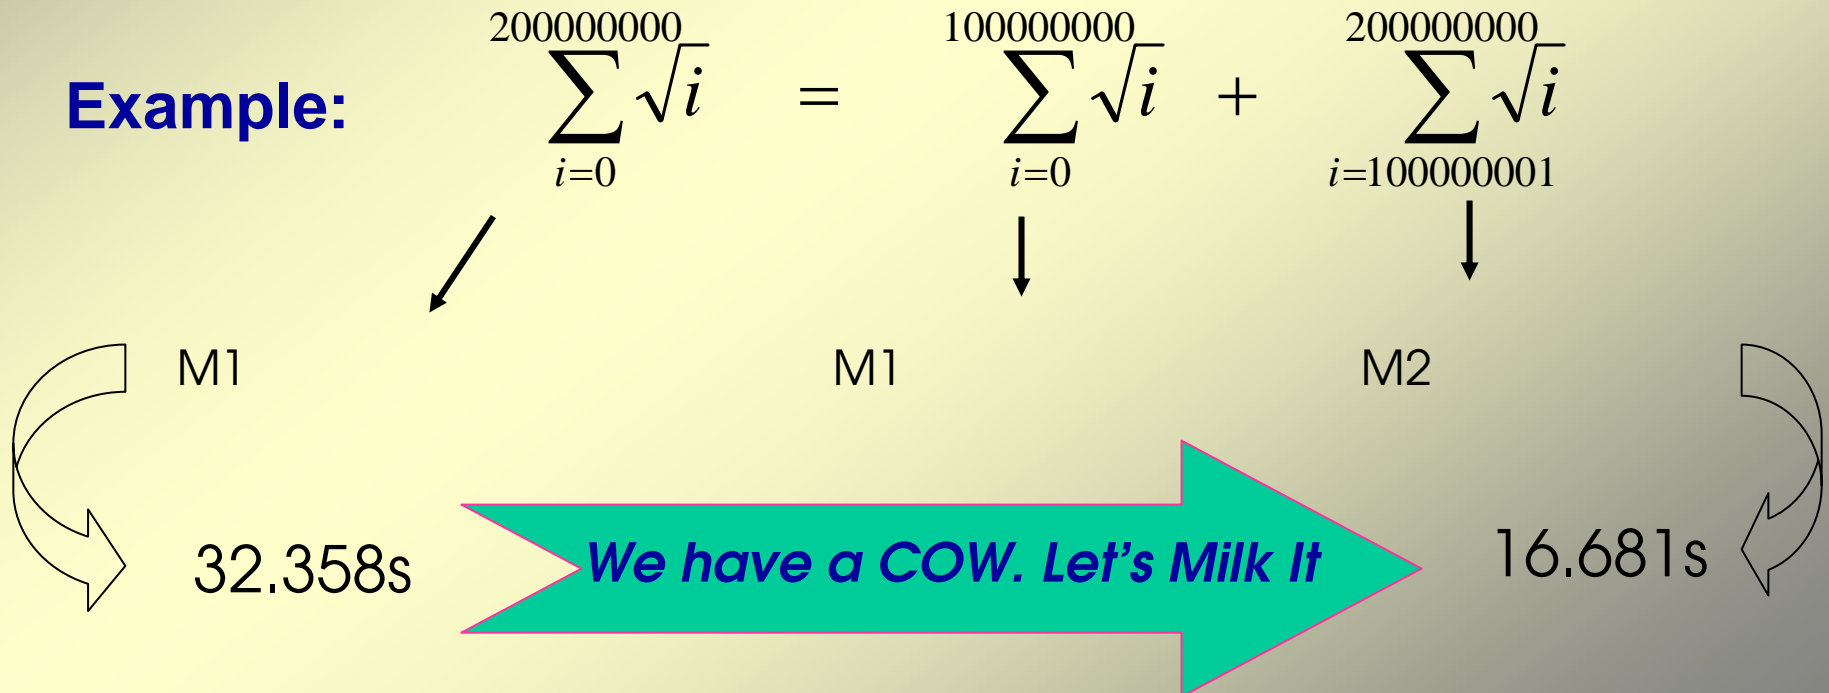
- One (typically) server node, one or more client nodes connected via Ethernet (or some other N/W)
- Built of commodity hardware
  - PC running Linux
  - Standard Ethernet adapters & switches
- No custom hardware components
- Trivially reproducible
- Uses Message Passing Interface (MPI) and Parallel Virtual Machine (PVM)
- Clients are in the private N/W
- It is NOT a new software package or N/W topology or the latest kernel

***Beowulf behaves like a single machine rather than many workstations***

**Eg.**

- We have Network File System (NFS) and NIS server  
Data is shared by all machines in the N/W
- We have machines of same configuration (not a significant point but eases the COW implementation)  
we can run the same code on these machines
- Use RSH (remote shell) to do the job. Don't forget to put them in background

**Example:**





## w.r.t SimSET

- Allocation of activity is concurrent
  - can be run on different machines
  - server doing the pre-processing to distribute the load
  - use bottom-up approach for load balancing
- No. of decays per cell/ grouped cells
  - use PRNGs with sufficiently longer period
  - search for PRNGs to be run on threads
- Cost analysis for installing/ maintain a Beowulf cluster
  - what hardware acceleration is possible/ feasible
  - TBD

***Beowulf gives better price/ performance ratio & is dedicated for parallel processing. Suitable for specific applications → SimSET !!!***

## DO's before clustering

- Determine the concurrent parts of the body
  - Number Crunching
  - I/O
- Estimate parallel efficiency
  - Good Old Trail & Error
  - Is it CPU limited/ Hard-Disk limited
  - Gut Feeling
  - A of DMA
- Describing the concurrent parts of the algorithm
  - Implicit (User)
  - Explicit (Compiler)

***A Compute bound problem may work better with few faster CPUs!  
An I/O bound problem may work better with more slow CPUs!!!***

# Resources/ Links

Beowulf Homepage <http://www.beowulf.org>

Extreme Linux <http://www.extremelinux.org>

The latest version of the Beowulf HOWTO <http://www.sci.usq.edu.au/staff/jacek/beowulf>

Building a Beowulf System <http://www.cacr.caltech.edu/beowulf/tutorial/building.html>

Linux Parallel Processing HOWTO <http://yara.ecn.purdue.edu/~{pplinux/PPHOWTO/pphowto.Html>

Cluster computing review <http://www.npac.syr.edu/techreports/hypertext/sccs-748/cluster-review.html>

Computing Resources for Macromolecular Crystallography at DND-CAT

<http://tomato.dnd.aps.anl.gov/macrocryst/resources.html>

Frequently Asked Questions, Tips, and Hints (IIT) <http://www.iit.edu/~gigawulf/faq.html>

Home Supercomputing with Linux <http://www.cris.com/~rjbono/html/pondermatic.html>

Useful Links <http://shspc1.chem.psu.edu/pa/int/linuxcluster.html>

Papers <http://www.beowulf.org/papers/HPDC96/hpdc96.html>

<http://www.beowulf.org/papers/AA97/aa97.ps>

<http://www.beowulf.org/papers/IPPS96/ipps96.html>

Price watch for hardware components <http://www.pricewatch.com/> , <http://computers.cnet.com/>

Why IBM is doing Linux Clusters [http://www-1.ibm.com/servers/eserver/clusters/whitepapers/linux\\_wp.html](http://www-1.ibm.com/servers/eserver/clusters/whitepapers/linux_wp.html)

GE CRD Whitney Information Services eBook Collection <http://www.netlibrary.com/index.asp>

How to build a Beowulf system

Software Updates, Archives, Documentation on Beowulf: <http://www.beowulf-underground.org/>

**Now Beowulf is hosted by <http://www.scyld.com/>**

**We don't offer full support for your Linux clustering requirements!!**

# IBM Linux Cluster

<http://www.cineca.it/HPSystems/Resources/LinuxCluster/>



Peak performance: 145 Gflops/s

Processor: PIII 1.1333GHz

No. : 128

Compilers: GCC C, C++, F90

OS: Linux

## Acknowledgements

Jack Radajewski & Douglas Eadline  
helpdesk  
and  
WWW

80% of this presentation was prepared using CAP Tools  
and is 100% recyclable

***No Surprise if you see  
“building super computers at home for dummy’s”***